

Jin Hong Moon

jaydenmoon7@gmail.com | jinhongmoon.com | github.com/jmo-on | US Permanent Resident

EDUCATION

Johns Hopkins University

Expected Dec 2026

B.S./M.S. in Computer Science, Cumulative GPA: 3.88

Baltimore, MD

Relevant Courses: Machine Learning, Deep Learning, Operating Systems, Parallel Programming, Algorithms, System Design, Object-oriented Design, Databases, Linear Algebra, Probability, Statistics

Teaching Assistant: Data Structures (Java), Intermediate Programming (C/C++)

EXPERIENCE

Software Engineering Intern

May 2026 – August 2026

Google

Sunnyvale, CA

- Designed and shipped a C++ caching layer eliminating redundant Gmail body-hash computations at 1.53M QPS
- Optimized cache memory access for CPU cache locality and safely reused hashes across MIME message mutations
- Validated zero cryptographic mismatches across billions of operations in production; bypassed 945K QPS, reduced hashing latency by 52.5%, and saved ~454 continuous CPU cores

Machine Learning Researcher

Jan 2026 – Present

Johns Hopkins University

Baltimore, MD

- Designed an energy-based vision transformer that jointly refines feature and attention maps via energy minimization, improving 1.4 mIoU over a baseline with near-zero added parameters
- Built a cross-layer feedback system that propagates converged attention states between transformer layers

Software Engineering Intern

June 2025 – August 2025

Bloomberg

New York City, NY

- Designed a distributed, scalable data pipeline for PORT<GO> time-series data layer, developing a 20+ threaded C++ producer, Python consumer, and a 4-partition Confluent Kafka topic
- Optimized read query performance of a C++ high-throughput, low-latency infrastructure processing 4B+ daily queries at 25K requests/sec, reducing average latency by 32.6%, significantly accelerating PORT<GO> analytics
- Built a real-time time-series analytics layer using ANSI SQL, Apache Trino, and Superset

Software Engineering Intern

June 2024 – November 2024

Claudius Legal Intelligence

Remote

- Developed 5+ API endpoints integrated with jQuery and AJAX frontends for seamless pagination and filtering
- Initiated AI legal assistance by implementing locally fine-tuned BERT and BART models with PEFT LoRA in PyTorch, generating trending article topics and recommendations

PROJECTS

Summit - Job Application Tracker | summit-jobs.com

February 2025 – May 2025

- Created a distributed task queue using Celery and Redis in Django REST Framework to parse job-related content from Gmail via OpenAI API with 92.6% accuracy, enabling automated job tracking from users' inboxes
- Improved a lightweight DistilBERT-based classifier to filter non-job-related emails before parsing with accuracy of 96.5%, significantly reducing token usage and improving throughput to 5K tokens/sec
- Built an OpenAI-powered chatbot with semantic search over 16K job postings by leveraging vector embeddings generated via text-embedding-3-small and stored in Supabase's pgvector, delivering personalized recommendations

Noori AI - Medical Translation | nooriai.com

February 2025 – May 2025

- Led development of a HIPAA-compliant sub-500ms medical interpreter application across 3 languages, reducing interpretation costs by over 90% for healthcare institutions serving patients with Limited English Proficiency
- Achieved 37.0 BLEU score on a MarianNMT model trained on OPUS and web-scraped medical data, preprocessed with custom tokenizers to produce aligned sentence pairs, under the mentorship of Dr. Philipp Koehn

TECHNICAL SKILLS

Languages: Python, C, C++, Go, Java, SQL

Frameworks: Flask, Django, FastAPI, Spring Boot, Node.js, React.js, Next.js

Developer Tools: Docker, Kubernetes, AWS, GCP, Supabase, Apache Kafka, Apache Spark, Apache Trino

Libraries: PyTorch, TensorFlow, Scikit-learn, HuggingFace, Numpy, Pandas, Celery, gMock