

Jin Hong Moon

jaydenmoon7@gmail.com | jinhongmoon.com | github.com/jmo-on | US Permanent Resident

EDUCATION

Johns Hopkins University

Expected Dec 2026

Master of Science in Engineering in Computer Science, Cumulative GPA: 3.93

Baltimore, MD

Bachelor of Science in Computer Science

Relevant Courses: Machine Learning, Deep Learning, Artificial Intelligence, Object-oriented Engineering Algorithms, Computer Systems, Databases, Linear Algebra, Probability, Statistics, Graphics

Teaching Assistant: Data Structures (Java), Intermediate Programming (C/C++)

EXPERIENCE

Software Engineering Intern

June 2025 – August 2025

Bloomberg

New York City, NY

- Designed a distributed, scalable data pipeline for PORT(GO) time-series data layer, developing a 20+ threaded C++ producer, Python consumer, and a 4-partition Confluent Kafka topic
- Optimized read query performance of a C++ high-throughput, low-latency infrastructure processing 4B+ daily requests at 25K requests/sec, reducing average latency by 32.6%, significantly accelerating PORT(GO) analytics
- Built a time-series analytics layer using Apache Trino and Superset, querying across 5 internal databases and APIs with ANSI SQL to provide real-time insights and 20+ years of historical reference data
- Owned end-to-end development, managing Jira tickets for Agile sprints and ensuring quality through code reviews

Software Engineering Intern

June 2024 – November 2024

Claudius Legal Intelligence

Remote

- Developed 5+ API endpoints for paper submission, rejection, and folder creation, powered by Stripe payments, integrated with 5+ frontend pages via jQuery and AJAX for pagination and filtering
- Initiated AI legal assistance by implementing locally fine-tuned BERT and BART models with PEFT LoRA in PyTorch, generating trending article topics and recommendations
- Deployed models on Google Cloud Functions to process data from Google Cloud Storage and established CI/CD pipeline through Google Cloud Build and Google App Engine for Dockerized deployments

PROJECTS

Summit - Job Application Tracker | summit-jobs.com

February 2025 – May 2025

- Created a distributed task queue using Celery and Redis in Django REST Framework to parse job-related content from Gmail emails via OpenAI API with 92.6% accuracy, enabling automated job tracking from users' inboxes
- Improved a lightweight DistilBERT-based classifier to filter non-job-related emails before parsing with accuracy of 96.5%, significantly reducing token usage and improving latency to 5K tokens/sec
- Automated Gmail Watcher renewal and token refresh through GCP Pub/Sub and Celery Beat, ensuring continuous Gmail API subscriptions and an uninterrupted inbox-to-dashboard pipeline
- Built an OpenAI-powered chatbot with semantic search over 16K job postings by leveraging vector embeddings generated via text-embedding-3-small and stored in Supabase's pgvector, delivering personalized recommendations
- Spearheaded team coordination by performing Agile sprints, peer-reviewed PRs, and comprehensive documentation

Noori AI - Medical Translation | nooriai.com

February 2025 – May 2025

- Led development of a HIPAA-compliant medical interpreter application across 3 languages, reducing interpretation costs by over 90% for healthcare institutions serving patients with Limited English Proficiency
- Achieved 37.0 BLEU score on a MarianNMT model trained on OPUS and web-scraped medical data, preprocessed with custom tokenizers to produce aligned sentence pairs, under the mentorship of Dr. Philipp Koehn
- Engineered a fully offline, sub-500ms speech-to-text translation pipeline using PyAudio and VOSK integrated with the MarianNMT model, ensuring instant communication across clinical environments

TECHNICAL SKILLS

Languages: Python, C, C++, Go, Java, SQL

Frameworks: Flask, Django, FastAPI, Spring Boot, Supabase, Node.js, React.js, Next.js

Developer Tools: Docker, Kubernetes, CMake, AWS, GCP, Apache Kafka, Apache Spark, Apache Trino

Libraries: Pytorch, Tensorflow, Scikit-learn, HuggingFace, Numpy, Pandas, Celery, gMock